Before you turn this problem in, make sure everything runs as expected. First, **restart the kernel** (in the menubar, select Kernel→Restart) and then **run all cells** (in the menubar, select Cell→Run All).

Make sure you fill in any place that says `YOUR CODE HERE` or "YOUR ANSWER HERE", as well as your name and collaborators below:

```
In [1]:  NAME = ""
         COLLABORATORS = ""
```

Problem description

In the last assignment, we created fully connected neural networks to solve the task of classifying whether a ship is present in a satellite photo.

This assignment will address the same task, but using Convolutional Neural Network layers.

Goal:

In this notebook, you will need to create a model in `TensorFlow/Keras` *using Convolutional layers* to classify satellite photos.

- The features are images: 3 dimensional collection of pixels
    - 2 spatial dimensions
    - 1 dimension with 3 features for different parts of the color spectrum: Red, Green, Blue
- The labels are either 1 (ship is present) or 0 (ship is not present)

Unlike the model we created in the previous assignment (using Dense/Fully Connected layers) we will retain *all three color channels* of the data rather than converting it to one channel gray scale.

There are two notebook files in this assignment:

- The one you are viewing now: First and only notebook you need to work on.
    - Train your models here
    - There are cells that will save your models to a file
- **Model_test.ipynb**:
    - PLEASE IGNORE

Learning objectives

- Learn how to construct Neural Networks in a Keras Sequential model that uses Convolutional layer types.
- Appreciate how layer choices impact number of weights

Import modules

```
In [2]:  ## Standard imports
         import numpy as np
         import pandas as pd
         import matplotlib.pyplot as plt

         import sklearn

         import os
         import math

         %matplotlib inline

         ## Import tensorflow
         import tensorflow as tf
         from tensorflow.keras.utils import plot_model

         print("Running TensorFlow version ",tf.__version__)

         # Parse tensorflow version
         import re

         version_match = re.match("([0-9]+)\.([0-9]+)", tf.__version__)
         tf_major, tf_minor = int(version_match.group(1)) , int(version_match.group(2))
         print("Version {v:d}, minor {m:d}".format(v=tf_major, m=tf_minor) )
```

```
Running TensorFlow version  2.4.1
Version 2, minor 4
```

API for students

We have defined some utility routines in a file `helper.py`. There is a class named `Helper` in it.

This will simplify problem solving

More importantly: it adds structure to your submission so that it may be easily graded

```
helper = helper.Helper()
```

- getData: Get a collection of labeled images, used as follows

data, labels = helper.getData()

```
In [3]:   # Load the helper module
          from IPython.core.interactiveshell import InteractiveShell
          InteractiveShell.ast_node_interactivity = "all"

          # Reload all modules imported with %aimport
          %reload_ext autoreload
          %autoreload 1

          # Import nn_helper module
          import helper
          %aimport helper
```

You may pass an optional `data_dir` argument to the constructor `helper.Helper`

- if you have your data directory located somewhere other than the default location
- no need to provide an argument otherwise

For example:

```
helper = helper.Helper(data_dir="./MyDataDir")
```

```
In [4]:   helper = helper.Helper()
```

Get the data

The first step in our Recipe is Get the Data.

We have provided a utility method `getData` to simplify this for you

```
In [ ]:   # Get the data
          data, labels = helper.getData()
          n_samples, width, height, channel = data.shape

          print("Data shape: ", data.shape)
          print("Labels shape: ", labels.shape)
          print("Label values: ", np.unique(labels))
```

We will shuffle the examples before doing anything else.

This is usually a good idea

- Many datasets are naturally arranged in a *non-random* order, e.g., examples with the sample label grouped together
- You want to make sure that, when you split the examples into training and test examples, each split has a similar distribution of examples

```
In [ ]:   # Shuffle the data
          data, labels = sklearn.utils.shuffle(data, labels, random_state=42)
```

Have a look at the data

We will not go through all steps in the Recipe, nor in depth.

But here's a peek

```
In [ ]:   # Visualize the data samples
          helper.showData(data[:25], labels[:25])
```

Have look at the data: Examine the image/label pairs

Rather than viewing the examples in random order, let's group them by label.

Perhaps we will learn something about the characteristics of images that contain ships.

We have loaded and shuffled our dataset, now we will take a look at image/label pairs.

Feel free to explore the data using your own ideas and techniques.

```
In [ ]:    # Inspect some data (images)
           num_each_label = 10

           for lab in np.unique(labels):
               # Fetch images with different labels
               X_lab, y_lab = data[ labels == lab ], labels[ labels == lab]
               # Display images
               fig = helper.showData( X_lab[:num_each_label], [ str(label) for label in y_l
           ab[:num_each_label] ], cmap="gray")
               _ = fig.suptitle("Label: "+  str(lab), fontsize=14)
               print("\n\n")
```

Make sure the features are in the range [0,1]

Warm up exercise: When we want to train on image data, the first thing we usually need to do is scaling.

Since the feature values in our image data are between 0 and 255, to make them between 0 and 1, we need to divide them by 255.

We also need to consider how to represent our target values

- If there are more than 2 possible target values, One Hot Encoding may be appropriate
    - **Hint**: Lookup `tf.keras.utils.to_categorical`
- If there are only 2 possible targets with values 0 and 1 we can use these targets without further encoding

Question

- Set variable X to be our **color** (3 channels) examples (`data`), but with values in the range [0,1]
- Set variable y to be the representation of our target values

In [ ]:
```
# Scale the data
# Assign values for X, y
#   X: the array of features
#   y: the array of labels
# The length of X and y should be identical and equal to the length of data.
from tensorflow.keras.utils import to_categorical
X, y = np.array([]), np.array([])

# YOUR CODE HERE
raise NotImplementedError()
```

In [ ]:
```
# Check if your solution is right

assert X.shape == (4000, 80, 80, 3)
assert ( (y.shape == (4000,)) or (y.shape == (4000,1)) )
```

Split data into training data and testing data

To train and evaluate a model, we need to split the original dataset into a training subset (in-sample) and a test subset (out of sample).

We will do this for you in the cell below.

DO NOT shuffle the data until after we have performed the split into train/test sets

- We want everyone to have the **identical** test set for grading
- Do not change this cell

```
In [ ]:  # Split data into train and test
         from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10, rando
         m_state=42)

         # Save X_train, X_test, y_train, y_test for final testing
         if not os.path.exists('./data'):
             os.mkdir('./data')
         np.savez_compressed('./data/train_test_data.npz', X_train=X_train, X_test=X_tes
         t, y_train=y_train, y_test=y_test)
```

Create a simple CNN model

Question: Create a Keras Sequential model:

- With a single Convolutional Layer providing the number of features given in variable `num_conv_features1` indicated below.
  - You may choose your own kernel size
  - Use same padding
- Feeding a head layer implementing Classification

Please name your Convolutional layer "CNN_1" and your Dense layer (head layer) "dense_head"

Set variable `model0` to be a Keras `Sequential` model object that implements your model.

Hints:

1. The `input shape` argument of the first layer should be the shape of a single example, which should be 3-dimensional. We don't need to flatten the data before feeding the Convolutional layer.
2. What is the shape of the output of the Convolutional Layer ? What should be the shape of the input to the Classification head ?
   - You may want to flatten the output of the Convolutional layer before feeding the Classification head.

3. The number of units in your head layer

   - Depends on how you represented the target
   - It should be equal to the final dimension of y

4. Activation function for the head layer: Since this is a classification problem

   - Use `sigmoid` if your target's final dimension equals 1
   - Use `softmax` if your target's final dimension is greater than 1

Think about the number of parameters in this model and whether it is likely to be prone to overfitting. If you feel the need:

- A Dropout layer maybe helpful to prevent overfitting and accelerate your training process

```python
In [ ]:  # Set model0 equal to a Keras Sequential model
         model0 = None

         num_conv_features1 = 32

         # YOUR CODE HERE
         raise NotImplementedError()

         model0.summary()
```

```python
In [ ]:  # Plot your model
         plot_model(model0)
```

Train model

Question:

Now that you have built your first model, you will compile and train it. The requirements are as follows:

- Split the **training** examples `X_train, y_train` again!

    - 80% will be used for training the model
    - 20% will be used as validation (out of sample) examples
    - Use `train_test_split()` from `sklearn` to perform this split
        - Set the `random_state` parameter of `train_test_split()` to be 42

- Loss function:

    - `binary_crossentropy` if your target is one-dimensional
    - `categorical_crossentropy` if your target is One Hot Encoded
- Metric: "accuracy"

- Use the value in variable `max_epochs` as the number of epochs for training

```
In [ ]:  model_name0 = "CNN + Head"
         max_epochs = 10

         # YOUR CODE HERE
         raise NotImplementedError()
```

How many weights in the model ?

Question:

Calculate the number of parameters in your model.

Set variable `num_parameters0` to be equal to the number of parameters in your model.

Hint: The model object may have a method to help you ! Remember that Jupyter can help you find the methods that an object implements.

```
In [ ]:  # Set num_parameters2 equal to the number of weights in the model
         num_parameters0 = None

         # YOUR CODE HERE
         raise NotImplementedError()

         print("Parameters number in model0: ", num_parameters0)
```

Evaluate the model

Question:

We have trained our model. We now need to evaluate the model using the test dataset created in an earlier cell.

Please store the model score in a variable named `score0`.

Hint: The model object has a method `evaluate`. Use that to compute the score.

```
In [ ]: score0 = []

        # YOUR CODE HERE
        raise NotImplementedError()

        print("{n:s}: Test loss: {l:3.2f} / Test accuracy: {a:3.2f}".format(n=model_name
        0, l=score0[0], a=score0[1]))
```

Your test accuracy should be around 0.9

Save the trained model0 and history for submission

Your fitted model can be saved for later use

- In general: so you can resume training at a later time
- In particular: to allow us to grade it !

Execute the following cell to save your model, which you will submit to us for grading.

```
In [ ]: helper.saveModel(model0, model_name0)
        helper.saveHistory(history0, model_name0)
```

```
In [ ]:   ## Restore the model (make sure that it works)
          model_loaded = helper.loadModel(model_name0)
          score_loaded = model_loaded.evaluate(X_test, y_test, verbose=0)

          assert score_loaded[0] == score0[0] and score_loaded[1] == score0[1]
```

Create a model with 4 Convolutional layers

Question:

We will now create a model with more Convolutional layers.

- Use **4** Convolutional layers.

  - You may choose your own kernel size
  - Use full padding

  - ReLU activation functions for the Convolutional layers

  - Each of the first two Convolutional layers should have the number of features given in variable `num_conv_features1` indicated below

    - Please name thee layers "CNN_1" and "CNN_2"
  - Each of the last two Convolutional layers should have the number of features given in variable `num_conv_features2` indicated below
    - Please name thee layers "CNN_3" and "CNN_4"

- Insert a `MaxPooling` layer after every two Convolutional layers (e.g., after

```
In [ ]:  # Set model1 equal to a Keras Sequential model
         model1 = None
         num_conv_features1 = 32
         num_conv_features2 = 64

         # YOUR CODE HERE
         raise NotImplementedError()

         model1.summary()
```

```
In [ ]:   # Plot your model
          plot_model(model1)
```

Train model

Question:

Train your new model following the same instructions as given for training the first model.

- Use the same datasets for training and validation as in your first model
- Use the same Loss function and metrics as in your first model
- **Except**: Save your training results in a variable named `history1`

```
In [ ]:   # Train the model using the API
          model_name1 = "4CNNs + Head"

          # YOUR CODE HERE
          raise NotImplementedError()
```

How many weights in this model ?

Question: Calculate the number of parameters in your new model.

Set variable `num_parameters1` to be equal to the number of parameters in your

```
In [ ]:  # Set num_parameters3 equal to the number of weights in the model
         num_parameters1 = None

         # YOUR CODE HERE
         raise NotImplementedError()

         print('Parameters number in model1:', num_parameters1)
```

Evaluate the model

Evaluate your new model following the same instructions as given for evaluating the first model.

- **Except**: store the model score in a variable named `score1`.

```
In [ ]:  score1 = []

         # YOUR CODE HERE
         raise NotImplementedError()

         print("{n:s}: Test loss: {l:3.2f} / Test accuracy: {a:3.2f}".format(n=model_name
         1, l=score1[0], a=score1[1]))
```

Is your test accuracy higher than before ?